



ROBUSTNESS





Z SCORE RECAP



Z SCORE

-
- z-score describes the location of the raw score
 - in terms of distance from the mean, measured in standard deviations
 - Gives us information about the location of that score relative to the “average” deviation of all scores
 - A z-score is the number of standard deviations a score is above or below the mean of the scores in a distribution.
 - A raw score is a regular score before it has been converted into a Z score
 - Raw scores on very different variables can be converted into Z scores and directly compared

Z SCORE

Raw scores on very different variables can be converted into Z scores and directly compared

What does a z-score tell us

- Mean of zero
- Zero distance from the mean
- Standard deviation of 1
- The z-score has two parts
 - The number
 - The sign

CONVERTING SCORES INTO Z SCORES AND BACK

-
- Now we will consider several examples that illustrate the use of the z-score formula. Suppose that we know about a population of a particular breed of cats having weights that are normally distributed. Furthermore, suppose we know that the mean of distribution is 10 pounds and the standard deviation is 2 pounds. Consider the following questions:
 - What is the z-score for 13 pounds?
 - What is the z-score for 6 pounds?
 - How many pounds corresponds to a z-score of 1.25?

CONVERTING SCORES INTO Z SCORES AND BACK

- What is the z-score for 13 pounds?
 - $(13 - 10)/2 = 1.5$
 - This means that 13 is one and a half standard deviations above the mean.
- What is the z-score for 6 pounds?
 - $(6 - 10)/2 = -2$
 - The interpretation of this is that 6 is two standard deviations below the mean.
- How many pounds corresponds to a z-score of 1.25?
 - $1.25 = (x - 10)/2$
 - $2.5 = (x - 10)$
 - $12.5 = x$
 - And so we see that 12.5 pounds corresponds to a z-score of 1.25.

CONVERTING ACROSS MEASURES

- We can convert from one metric to another.
- National results for the SAT test show that for college-bound seniors the average combined SAT Writing, Math and Verbal score is 1500 and the standard deviation is 250.
- National results for the ACT test show that for college-bound seniors the average composite ACT score is 20.8 and the standard deviation is 4.8.
 - Your SAT score: 1860.
 - Your neighbor's ACT: 29.
- Who did better on their respective test, and how would each have done on the other test?

CONVERTING ACROSS MEASURES

- Who did better on their respective test?
- SAT into Z
 - $1860 - 1500 = 360$
 - $360 / 250 = 1.44$
- ACT into Z
 - $29 - 20.8 = 8.2$
 - $8.2 / 4.8 = 1.708$

CONVERTING ACROSS MEASURES

-
- How would each have done on the other test?
 - ZSAT into ACT
 - $1.44 * 4.8 = 6.912$
 - $6.912 + 20.8 = 27.712$
 - ZACT into SAT
 - $1.708 * 250 = 427$
 - $427 + 1500 = 1927$

COMMENTS ON STANDARDIZATION

- Standardization of variables
 - Transformation – ratio variable
 - Features of measurement scales
 - Some have built in properties
 - Percentiles (not equal interval!)
 - Standardized scale (z score scale)

PERCENTILES




- Percentiles (percentile rank)
- Score value that cuts off a certain % of scores below it (cumulative percentile)
- Special percentile ranks
 - M: 50%
 - Q1: 25%
 - Q3: 75%
- Don't have equal intervals
 - Weak statistically



SKEWNESS

SKEWNESS

- Cowardly Skier



POSITIVE SKEW


- The distribution is said to be *right-skewed*, *right-tailed*, or *skewed to the right*,
 - *despite* the fact that the curve itself appears to be skewed or leaning to the left;
- The right tail is longer;
 - the mass of the distribution is concentrated on the left
- *right* instead refers to the right tail being drawn out and, often,
- A right-skewed distribution usually appears as a *left-leaning* curve.

NEGATIVE SKEW FACTS

- Also called: *left-skewed*, *left-tailed*, or *skewed to the left*
- The left tail is longer;
- the mass of the distribution is concentrated on the right of the figure.
- *left* instead refers to the left tail being drawn out
- A left-skewed distribution usually appears as a *right-leaning* curve.


RELATIONSHIP BETWEEN MEAN, MEDIAN, AND MODE

- Symmetric unimodal distribution
 - Mean=median=mode




WORKED EXAMPLE

- Where do students go to school?
 - Although 80.4% of first-time first-year students attended college in the state in which they lived, this percent varied considerably over the states.
 - Here is a stemplot of the percent of first-year students in each of the 50 states who were from the state where they enrolled.
 - The stems are 10s and the leaves are 1s.
 - The stems have been split in the plot.




WORKED EXAMPLE

- Where do students go to school?
 - Although 80.4% of first-time first-year students attended college in the state in which they lived, this percent varied considerably over the states.
 - Here is a stemplot of the percent of first-year students in each of the 50 states who were from the state where they enrolled.
 - The stems are 10s and the leaves are 1s.
 - The stems have been split in the plot.




WORKED EXAMPLE

- Where do students go to school?
 - Although 80.4% of first-time first-year students attended college in the state in which they lived, this percent varied considerably over the states.
 - Here is a stemplot of the percent of first-year students in each of the 50 states who were from the state where they enrolled.
 - The stems are 10s and the leaves are 1s.
 - The stems have been split in the plot.
 - Questions
 - Where is the median?
 - The shape of the distribution is ____ skewed
 - The state with the smallest percent of first-year students enrolled in the state has what %
 - The state with the largest percent has %?




■ Questions

- Where is the median?
- The shape of the distribution is _____ skewed
- The state with the smallest percent of first-year students enrolled in the state has v
- The state with the largest percent has %?



■ Questions

- Where is the median?
 - 76
- The shape of the distribution is left skewed
- The state with the smallest percent of first-year students enrolled in the state has v
 - 34% enrolled.
- The state with the largest percent has %?
 - 94%





Outliers



THE STORY OF SUCCESS

MALCOLM
GLADWELL

#1 bestselling author of *The Tipping Point* and *Blink*

ROBUST STATISTICS

OUTLIERS

The Story of Success

MALCOLM GLADWELL

INTRODUCTION

The Roseto Mystery

“THESE PEOPLE WERE DYING
OF OLD AGE. THAT’S IT.”

out·li·er \-,lī(-ə)r\ *noun*

1: something that is situated away from or classed differently from a main or related body

2: a statistical observation that is markedly different in value from the others of the sample

ROBUST STATISTICS AND OUTLIERS

INTRODUCTION

- out-li-er noun
- 1: something that is situated away from or classed differently from a main or related body
- 2: a statistical observation that is markedly different in value from the others of the sample
- “An important kind of deviation is an **outlier**, an individual value that falls outside the overall pattern.”

The Roseto Mystery

“THESE PEOPLE WERE DYING
OF OLD AGE. THAT’S IT.”


out·li·er \-,lī(-ə)r\ *noun*

1: something that is situated away from or classed differently from a main or related body

2: a statistical observation that is markedly different in value from the others of the sample

ROBUSTNESS

- Robust Statistics are less sensitive to outliers
- Most statistics, like
 - means, standard deviations, and correlations, and
 - every statistic based on these, are highly sensitive to outliers.



SO WHAT DO WE
DO WITH
OUTLIERS?

Do we remove
them?

Do we ignore them?

EXAMPLE

- The data shows Tukey's scores for the last 5 math tests: 88, 90, 55, 94, and 89.
- Identify the outlier in the data set.
- Then determine how the outlier affects the mean, median, and mode of the data

CALCULATIONS: MEAN

With Outlier

- Mean: $55+88+89+90+94 = 416$
 - $416 / 5 = 83.2$
 - The mean is 83.2

Without

- Mean: $88+89+90+94 = 361$
 - $361 / 4 = 90.25$
 - The mean is 90.25.

CALCULATIONS: MEDIAN

With Outlier

- Median: 55,88,89,90,94
 - The median is 89

Without

- Median: 88,89,90,94
 - $(89+90) / 2$
 - The median is 89.5

SO WHAT DO WE DO WITH OUTLIERS?



- Do we remove them?
- Do we ignore them?
 - It's up to the researcher

SO WHAT DO WE DO WITH OUTLIERS?

- Do we remove them?
 - If you remove them, you need to document it
 - Why did you remove them?
 - Obvious error? Impossible number?
- Do we ignore them?
 - Easier, but it might alter our results



5- NUMBER SUMMARY TABLE

- Summarize a distribution
- Old (mean, standard deviation)
- Tukey's (lower extreme, lower hinge, median, upper hinge, upper extreme)
 - Commonly referred to (min, q1, m, q3, max)



5- NUMBER SUMMARY TABLE









- 5- number summary table
 - Summarize a distribution
 - Old (mean, standard deviation)
 - Tukey's (lower extreme, lower hinge, median, upper hinge, upper extreme)
 - Commonly referred to (min, q1, m, q3, max)
 - Graph of the 5-summary table called boxplot
 - (lots of variants on this plot, including
 - Beeswarm Boxplot




DIAGNOSING OUTLIERS

- IQR rule (IQR= Interquartile Range ($Q3-Q1$)
 - $Q3 + 1.5 \times IQR$
 - $Q1 - 1.5 \times IQR$
 - Creates the fences outside of which the items
- Other Rules
 - 2.5 Standard Deviations beyond the mean
 - More Complex Algorithms
 - [Illustration with Rshiny](#)

The Datasets



WORKED EXAMPLE IN R



```
library(dslabs)
library(dplyr)
library(ggplot2)

# These data are self-reported heights
data(reported_heights)
str(reported_heights)

# Here we create a new variable within the reported heights data
reported_heights <- reported_heights %>%
  mutate(original_heights = height, height = as.numeric(height))

# Here we exclude heights with missing data
```

The screenshot shows the RStudio interface with a script editor. The toolbar at the top includes icons for file operations, a search icon, a 'Source on Save' checkbox, and buttons for 'Run', 'Source', and 'Source on Save'. The script content is as follows:

R CODE FOR BEESWARM

```
if(!require(beeswarm)) install.packages(beeswarm)
data(breast)
beeswarm(time_survival ~ event_survival,
  data = breast, method = 'swarm', pch = 16,
  pwcol = as.numeric(ER), xlab = '',
  ylab = 'Follow-up time (months)',
  labels = c('Censored', 'Metastasis'))

boxplot(time_survival ~ event_survival, data = breast, add =
T, names = c("", ""), col="#0000ff22")
```

